

BLAKE3

one function, fast everywhere

Jack O'Connor (@oconnor663)
Jean-Philippe Aumasson (@veorq)
Samuel Neves (@sevenps)
Zooko Wilcox-O'Hearn (@zooko)

<https://blake3.io>

We present BLAKE3, an evolution of the hash BLAKE2 that is faster, simpler to use, and better suited to applications' needs. BLAKE3 supports an unbounded degree of parallelism, using a tree structure that scales up to any number of SIMD lanes and CPU cores. On Intel Kaby Lake, peak single-threaded throughput is $3\times$ that of BLAKE2b, $4\times$ that of SHA-2, and $8\times$ that of SHA-3, and it scales further using multiple threads. BLAKE3 is also efficient on smaller architectures and microcontrollers: throughput on the 32-bit ARM1176 core is $1.3\times$ that of BLAKE2s or SHA-256, and $3\times$ that of BLAKE2b, SHA-512, or SHA-3. Unlike BLAKE2 and SHA-2, which have incompatible variants better suited for different platforms, BLAKE3 is a single hash function, designed to be consistently fast across a wide variety of software platforms and use cases.

Contents

1	Introduction	3
2	Specification	4
2.1	Tree Structure	4
2.2	Compression Function	5
2.3	Modes	6
2.4	Chunk Chaining Values	6
2.5	Parent Node Chaining Values	7
2.6	Extendable Output	7
3	Performance	8
4	Security	10
4.1	Security Goals	10
4.2	Mode of Operation	11
4.3	Compression Function	12
4.4	Cryptanalysis	12
5	Implementation	12
5.1	Incremental Hashing	12
5.1.1	Chunk State	13
5.1.2	Chaining Value Stack	13
5.2	Multi-threading	15
5.3	SIMD	15
5.4	Memory Requirements	16
6	Applications	17
6.1	Pseudorandom Function and MAC	17
6.2	Key Derivation	17
6.3	Stateful Hash Object	18
6.4	Verified Streaming	18
6.5	Incremental Update	20
7	Rationales	20
7.1	Chunk Size	20
7.2	Word Size	21
7.3	Compression Function Feed Forward	21
7.4	Chunk Counter	22
7.5	Tree Fanout	22
7.6	Domain Flags	23
7.7	Key Derivation from Reused Key Material	24
A	IV Constants	27
B	Round Function	27

1 Introduction

Since its announcement in 2012, BLAKE2 [5] has seen widespread adoption, in large part because of its superior performance in software. BLAKE2b and BLAKE2s are included in OpenSSL and in the Python and Go standard libraries. BLAKE2b is also included as the `b2sum` utility in GNU Coreutils, as the `generichash` API in Libsodium, and as the underlying hash function for Argon2 [8], the winner of the Password Hashing Competition in 2015.

The biggest changes from BLAKE2 to BLAKE3 are:

- An **internal tree structure**.
- A compression function with **fewer rounds**.
- A **single hash function**, with no variants or flavors.
- In lieu of a parameter block, an API supporting **three domain-separated modes**: `hash(input)`, `keyed_hash(key, input)`, and `derive_key(context, key_material)`.
- The space formerly occupied by the parameter block is now used for the optional 256-bit key, so **keying is zero-cost**.
- BLAKE3 has built-in support for **extendable output**. Like BLAKE2X, but unlike SHA-3 or HKDF, extended output is parallelizable and seekable.

BLAKE3 thus eliminates the main drawback of BLAKE2, namely its number of incompatible variants; the original BLAKE2 paper described 64-bit BLAKE2b, 32-bit BLAKE2s, the parallel variants BLAKE2bp and BLAKE2sp, a framework for tree modes, and the BLAKE2X paper later added extendable output modes. None of these are compatible with each other, and choosing the right one for an application means understanding both the tradeoffs between them and also the state of language and library support. BLAKE2b, the most widely supported, is not the fastest on most platforms. BLAKE2bp and BLAKE2sp, with much higher peak throughput, are sparsely supported and rarely adopted. BLAKE3 instead is a single hash function, with no variants, designed to support all the use cases of BLAKE2, as well as new use cases like verified streaming (see §6.4).

BLAKE3 is also dramatically faster. For example, on an Intel Kaby Lake processor peak throughput on a single core is triple that of BLAKE2b, and BLAKE3 can scale further to any number of cores. On an ARM1176, throughput is $1.3\times$ that of BLAKE2s and again triple that of BLAKE2b. To achieve such speed, BLAKE3 splits its input into 1 KiB chunks and arranges those chunks as the leaves of a binary tree. This tree structure means that there is no limit to the parallelism that BLAKE3 can exploit, given enough input [1, 2]. The direct benefit of this parallelism is very high throughput on platforms with SIMD support, including all modern x86 processors. Another benefit of hashing chunks in parallel is that the implementation can use SIMD vectors of any width, regardless of the word size of the compression function. That leaves us free to use a compression function that is efficient on smaller architectures, without sacrificing peak throughput on x86-64.

The BLAKE3 compression function is closely based on that of BLAKE2s. BLAKE3 has the same 128-bit security level and 256-bit default output size. The round function is identical, along with the IV constants **and the message schedule**. **Thus, cryptanalysis of BLAKE2 applies directly to BLAKE3**. Based on that analysis and with the benefit of hindsight, we

believe that BLAKE2 is overly conservative in its number of rounds, and BLAKE3 reduces the number of rounds from 10 to 7 (see [3] for detailed rationale). BLAKE3 also changes the setup and finalization steps of the compression function to support the internal tree structure, more efficient keying, and extendable output.

2 Specification

2.1 Tree Structure

BLAKE3 splits its input into chunks of up to 1024 bytes and arranges those chunks as the leaves of a binary tree. The last chunk may be shorter, but not empty, unless the entire input is empty. If there is only one chunk, that chunk is the root node and only node of the tree. Otherwise, the chunks are assembled with parent nodes, each parent node having exactly two children. The structure of the tree is determined by two rules:

1. Left subtrees are full. Each left subtree is a complete binary tree, with all its chunks at the same depth, and a number of chunks that is a power of 2.
2. Left subtrees are big. Each left subtree contains a number of chunks greater than or equal to the number of chunks in its sibling right subtree.

In other words, given a message of $n > 1024$ bytes, the left subtree consists of the first

$$2^{10 + \lceil \log_2(\lfloor \frac{n-1}{1024} \rfloor) \rceil}$$

bytes, and the right subtree consists of the remainder.

For example, trees from 1 to 4 chunks have the structure shown in Figure 1.




Figure 1: Example tree structures, from 1 to 4 chunks.

The compression function is used to derive a chaining value from each chunk and parent node. The chaining value of the root node, encoded as 32 bytes in little-endian order, is the default-length BLAKE3 hash of the input. BLAKE3 supports input of any byte length $0 \leq \ell < 2^{64}$.

Lots of references to “default length”. Given that the output is not different between different sizes—in theory, BLAKE3’s output is an infinite-length string—it’s much less of an issue than in BLAKE2.

2.2 Compression Function

The compression function takes an 8-word chaining value, a 16-word message block, and a 4-word parameter, and it returns a new 16-word value. A word is 32 bits. The inputs to the compression function are:

- The input chaining value, $h_0 \dots h_7$.
- The message block, $m_0 \dots m_{15}$.
- A 64-bit counter, $t = t_0, t_1$, with t_0 the lower order word and t_1 the higher order word.
- The number of input bytes in the block, b .
- A set of domain separation bit flags, d .

The compression function initializes its 16-word internal state $v_0 \dots v_{15}$ as follows:

$$\begin{pmatrix} v_0 & v_1 & v_2 & v_3 \\ v_4 & v_5 & v_6 & v_7 \\ v_8 & v_9 & v_{10} & v_{11} \\ v_{12} & v_{13} & v_{14} & v_{15} \end{pmatrix} \leftarrow \begin{pmatrix} h_0 & h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 & h_7 \\ IV_0 & IV_1 & IV_2 & IV_3 \\ t_0 & t_1 & b & d \end{pmatrix}$$

The $IV_0 \dots IV_7$ constants are the same as in BLAKE2s, and they are reproduced in Appendix A.

The compression function applies a 7-round keyed permutation $v' = E(m, v)$ to the state $v_0 \dots v_{15}$, keyed by the message $m_0 \dots m_{15}$. **The keyed permutation here is identical to that of BLAKE2s, and is reproduced in Appendix B.**

The output of the compression function $h'_0 \dots h'_{15}$ is defined as:

$$\begin{array}{ll} h'_0 \leftarrow v'_0 \oplus v'_8 & h'_8 \leftarrow v'_8 \oplus h_0 \\ h'_1 \leftarrow v'_1 \oplus v'_9 & h'_9 \leftarrow v'_9 \oplus h_1 \\ h'_2 \leftarrow v'_2 \oplus v'_{10} & h'_{10} \leftarrow v'_{10} \oplus h_2 \\ h'_3 \leftarrow v'_3 \oplus v'_{11} & h'_{11} \leftarrow v'_{11} \oplus h_3 \\ h'_4 \leftarrow v'_4 \oplus v'_{12} & h'_{12} \leftarrow v'_{12} \oplus h_4 \\ h'_5 \leftarrow v'_5 \oplus v'_{13} & h'_{13} \leftarrow v'_{13} \oplus h_5 \\ h'_6 \leftarrow v'_6 \oplus v'_{14} & h'_{14} \leftarrow v'_{14} \oplus h_6 \\ h'_7 \leftarrow v'_7 \oplus v'_{15} & h'_{15} \leftarrow v'_{15} \oplus h_7. \end{array}$$

If we define v_l (resp. v'_l) and v_h (resp. v'_h) as the first and last 8-words of the input (resp. output) of $E(m, v)$ as elements of \mathbb{F}_2^{256} , the compression function may be represented as the affine transformation

$$\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} v'_l \\ v'_h \end{pmatrix} + \begin{pmatrix} 0 \\ v_l \end{pmatrix}.$$

The output of the compression function is often truncated to produce 256-bit chaining values.

The compression function input d is a bitfield, with each individual flag consisting of a power of 2. The value of d is the sum of all the flags defined for a given compression. Their names and values are given in Table 1.

Table 1: Admissible values for input d in the BLAKE3 compression function.

Flag name	Value
CHUNK_START	2^0
CHUNK_END	2^1
PARENT	2^2
ROOT	2^3
KEYED_HASH	2^4
DERIVE_KEY_CONTEXT	2^5
DERIVE_KEY_MATERIAL	2^6

2.3 Modes

BLAKE3 defines three domain-separated modes: `hash`, `keyed_hash`, and `derive_key`. The first two modes differ from each other in their key words $k_0 \dots k_7$ and in the additional flags they set for every call to the compression function:

- `hash`: $k_0 \dots k_7$ are the constants $IV_0 \dots IV_7$, and no additional flags are set.
- `keyed_hash`: $k_0 \dots k_7$ are parsed in little-endian order from the 32-byte key given by the caller, and the `KEYED_HASH` flag is set for every compression.

The third mode, `derive_key`, has two stages. First the context string is hashed, with $k_0 \dots k_7$ set to the constants $IV_0 \dots IV_7$, and the `DERIVE_KEY_CONTEXT` flag set for every compression. Then the key material is hashed, with $k_0 \dots k_7$ set to the first 8 output words of the first stage, and the `DERIVE_KEY_MATERIAL` flag set for every compression.

2.4 Chunk Chaining Values

Processing a chunk is structurally similar to the sequential hashing mode of BLAKE2. Each chunk of up to 1024 bytes is split into blocks of up to 64 bytes. The last block of the last chunk may be shorter, but not empty, unless the entire input is empty. The last block, if necessary, is padded with zeros to be 64 bytes.

Each block is parsed in little-endian order into message words $m_0 \dots m_{15}$ and compressed. The input chaining value $h_0 \dots h_7$ for the first block of each chunk is comprised of the key words $k_0 \dots k_7$. The input chaining value for subsequent blocks in each chunk is the output of the truncated compression function for the previous block.

The remaining compression function parameters are handled as follows (see also Figure 2 for an example):

- The counter t for each block is the chunk index, i.e., 0 for all blocks in the first chunk, 1 for all blocks in the second chunk, and so on.
- The block length b is the number of input bytes in each block, i.e., 64 for all full blocks except the last block of the last chunk, which may be short.
- The first block of each chunk sets the `CHUNK_START` flag (cf. Table 1), and the last block of each chunk sets the `CHUNK_END` flag. If a chunk contains only one block, that block

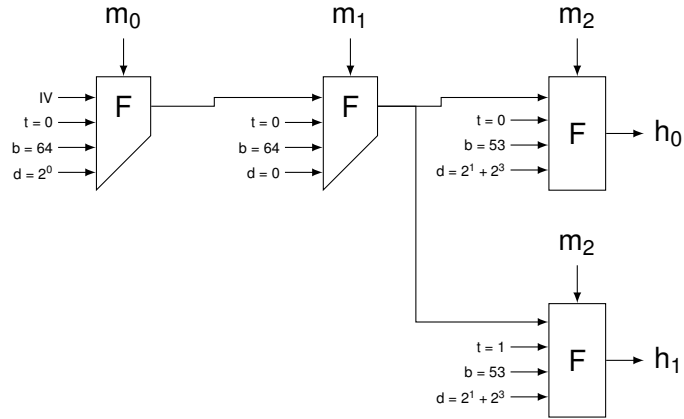


Figure 2: Example of compression function inputs when hashing a 181-byte input (m_0, m_1, m_2) into a 128-byte output (h_0, h_1). Trapezia indicate that the compression function output is truncated to 256 bits.

sets both `CHUNK_START` and `CHUNK_END`. If a chunk is the root of its tree, the last block of that chunk also sets the `ROOT` flag.

The output of the truncated compression function for the last block in a chunk is the chaining value of that chunk.

2.5 Parent Node Chaining Values

Each parent node has exactly two children, each either a chunk or another parent node. The chaining value of each parent node is given by a single call to the compression function. The input chaining value $h_0 \dots h_7$ is the key words $k_0 \dots k_7$. The message words $m_0 \dots m_7$ are the chaining value of the left child, and the message words $m_8 \dots m_{15}$ are the chaining value of the right child. The counter t for parent nodes is always 0. The number of bytes b for parent nodes is always 64. Parent nodes set the `PARENT` flag. If a parent node is the root of the tree, it also sets the `ROOT` flag. The output of the truncated compression function is the chaining value of the parent node.

2.6 Extendable Output

BLAKE3 can produce outputs of any byte length up to 2^{64} bytes. This is done by repeating the root compression—that is, the very last call to the compression function, which sets the `ROOT` flag—with incrementing values of the counter t . The results of these repeated root compressions are then concatenated to form the output.

When building the output, BLAKE3 uses the full output of the compression function (cf. §4.3). Each 16-word output is encoded as 64 bytes in little-endian order.

Observe that based on §2.4 and §2.5 above, the first root compression always uses the counter value $t = 0$. That is either because it is the last block of the only chunk, which has a chunk index of 0, or because it is a parent node. After the first root compression, as long as more output bytes are needed, t is incremented by 1, and the root compression is repeated on otherwise the same inputs. If the target output length is not a multiple of 64, the final compression output is truncated.

Because the repeated root compressions differ only in the value of t , the implementation can execute any number of them in parallel. The caller can also adjust t to seek to any point in the output stream.

Note that in contrast to BLAKE2 and BLAKE2X, BLAKE3 does not domain separate outputs of different lengths. Shorter outputs are prefixes of longer ones.

3 Performance




Figure 3: Throughput for a multi-threaded implementation of BLAKE3 at various input lengths on an Intel Kaby Lake processor.

Several factors contribute to the performance of BLAKE3, depending on the platform and the size of the input:

- The tree structure allows an implementation to compress multiple chunks in parallel using SIMD (cf. §5.3). With enough input, this can occupy SIMD vectors of any width.
- The tree structure allows an implementation to use multiple threads. With enough input, this can occupy any number of cores.
- The compression function uses fewer rounds than in BLAKE2.
- The compression function performs well on smaller architectures.




Figure 4: Throughput for a single-threaded implementation of BLAKE3 at various input lengths on an Intel Skylake-SP processor.

Figure 3 shows the throughput of a multi-threaded implementation of BLAKE3 on a typical laptop computer. This benchmark ran on an Intel Kaby Lake i5-8250U processor with 4 physical cores, supporting 256-bit AVX2 vector instructions, and with Turbo Boost disabled. In this setting, the implementation remains single-threaded for inputs up to 16 KiB. Parallel compression using SIMD begins with 8 KiB inputs using 128-bit vectors, and with 16 KiB inputs using 256-bit vectors. The implementation begins multi-threading with 32 KiB inputs, in this case causing a performance drop. (An implementation tuned for this specific platform would forgo multi-threading for inputs shorter than e.g. 128 KiB.) Peak throughput occurs at 4 MiB of input, at which point BLAKE3 is 5× faster than BLAKE2bp and BLAKE2sp and an order of magnitude faster than BLAKE2b, BLAKE2s, SHA-2, and SHA-3.

Figure 4 shows the throughput of a single-threaded implementation of BLAKE3 on modern server hardware. This benchmark ran on an AWS c5.metal instance with an Intel Skylake-SP 8175M processor, supporting AVX-512 vector instructions, and with Turbo Boost disabled. Here, BLAKE3 is the only algorithm able to take advantage of 512-bit vector arithmetic. (Note that the KangarooTwelve algorithm, not included in this benchmark, can also use 512-bit vectors.) The fixed tree structures of BLAKE2bp and BLAKE2sp limit those algorithms to 256-bit vectors.

Figure 5 shows the throughput of BLAKE3 on a smaller embedded platform. This benchmark ran on a Raspberry Pi Zero with a 32-bit ARM1176 processor, without multiple cores or SIMD support. Here, BLAKE2b, SHA-512, and SHA-3 perform relatively poorly, because




Figure 5: Throughput at various input lengths on an ARM1176 processor.

their compression functions require 64-bit arithmetic. BLAKE3 and BLAKE2s have similar performance profiles here, and the BLAKE3 compression function uses fewer rounds.

Figures 3, 4, and 5 also highlight that BLAKE3 performs well for short inputs. The fixed tree structures of BLAKE2bp and BLAKE2sp are costly when the input is short, because they always compress a parent node and a fixed number of leaves. The advantage of their fixed tree structures, however, comes at medium input lengths around 1–4 KiB, where BLAKE3 does not yet have enough chunks to operate in parallel. This is the regime where BLAKE2bp and BLAKE2sp pull ahead of BLAKE3 in figures 3 and 4.

4 Security

TODO

4.1 Security Goals

BLAKE3 targets 128-bit security for all of its security goals. That is, 128-bit security against (second-)preimage, collision, or differentiability attacks. Users who need 256-bit security can remain using BLAKE2.

The key length is nevertheless 256 bits, which is useful against multi-target or possible certification attacks.

On why 128-bit security is enough [3].

4.2 Mode of Operation

One important aspect of any hash mode, parallel or otherwise, is its *soundness*. Several works [7, 10–12] have studied some fairly general conditions on the requirements of the mode of operation of a hash function such that, when instantiated with an ideal compression function, block cipher, or permutation it remains indistinguishable from a random oracle. We adopt the requirements from Daemen et al. [11], which are:

- Subtree-freeness, i.e., no valid tree may be a subtree of another;
- Radical-decodability, i.e., preventing collisions by ambiguity of the tree shape;
- Message-completeness, i.e., the entire message can be recovered from the inputs to the primitive in question.

Subtree-freeness Subtree-freeness ensures that generalized length-extension attacks, like the ones that plagued Merkle-Damgård, cannot happen. To ensure subtree-freeness, BLAKE3 defines the `ROOT` flag, which is only set on the last compression output. Thus, for any valid tree, a valid subtree must have the same root node. There are two cases to consider here:

1. The `ROOT` flag is combined with `CHUNK_END`; in this case we are dealing with a single-chunk message, and the root nodes must be the same, as well as its parents, grandparents, etc, until we reach a `CHUNK_START` node. Because both subtrees must start with `CHUNK_START`, we conclude that both trees must be the same.
2. The `ROOT` flag is applied to a parent node. Here the root node necessarily has two parents, each of which must be equal, recursing until they hit the same set of leaves marked by the `CHUNK_END` flag.

Radical-decodability For radical-decodability, we can define the subset of final nodes. For each final node, one can define a function `radical()` that returns a CV position for any final node. Here we have, again, two cases:

1. The final node has the `CHUNK_END` flag set; here the radical is in the CV bits.
2. The final node is a parent node; here the CV comes in the message bits.

Message-completeness The entire input can be recovered from the message bits input to the chunks.

Then, by [11, Theorem 3], for a distinguisher \mathcal{D} of total complexity q we have

$$\text{Adv}_{\text{BLAKE3}}^{\text{diff}}(\mathcal{D}) \leq \frac{\binom{q}{2} + 1}{2^{256}} + \frac{\binom{q}{2}}{2^{512}} + \frac{q}{2^{256}}$$

for the 256-bit output. For the full-length output, we do not have a truncated permutation, but instead a feed-forward. We remark, however, that the truncation requirement in [11, Theorem 3] is in place purely to prevent the distinguisher from inverting an output; as such, the feed-forward serves the same purpose as truncation, and the result still follows.

Table 2: Best differential trail probabilities for increasing round numbers of the compression functions of BLAKE-256, BLAKE2s, and BLAKE3, respecting their constraints on inputs to the keyed permutation. Probabilities marked with * indicate symmetric differences were sought exclusively.

Function	0.5	1	1.5	2	2.5	3	3.5
BLAKE-256 c.f.	2^{-0}	2^{-0}	2^{-0}	2^{-1}	2^{-6}	2^{-7}	2^{-38}
BLAKE{2s,3} c.f.	2^{-0}	2^{-0}	2^{-0}	2^{-1}	2^{-32}	$2^{-88} \leq p < 2^{-48}$	-
BLAKE{-256,2s,3}	2^{-0}	2^{-1}	2^{-32}	$\geq 2^{-190*}$	-	-	-
BLAKE-256 c.f. (CV only)	2^{-0}	2^{-2}	2^{-12}	2^{-39}	-	-	-
BLAKE{2s,3} c.f. (CV only)	2^{-0}	2^{-7}	2^{-32}	$\geq 2^{-161*}$	-	-	-

4.3 Compression Function

4.4 Cryptanalysis

The rather aggressive round reduction from 10 to 7 rounds is based on existent cryptanalysis, along with novel research.

See [3].

- Cryptanalysis of BLAKE [4, 6, 9, 14, 18, 20, 21]
- Cryptanalysis of BLAKE2 [15–17]
- Table 2 shows that the initialization of BLAKE2/3, by restricting inputs, makes things harder for the attacker, with probabilities getting very low very quickly. As also observed in [16, §7].
- The boomerang attacks of, e.g., [6, 9, 17] exploit the high probabilities on the first few rounds to connect two trails and get through a few more rounds. This is a usual trick to get around the very quick lowering of probabilities in ARX constructs, along with differential-linear attacks. But such attacks do not present a threat to BLAKE3, considering its 128-bit security target, and input restrictions on the compression function input (i.e., the IV).

5 Implementation

5.1 Incremental Hashing

An incremental implementation of BLAKE3 has two major components: the state of the current chunk and a stack of subtree chaining values (the “CV stack”). The chunk state is structurally similar to an incremental implementation of the sequential mode of BLAKE2, and it will be familiar to implementers of other hash functions. The CV stack is less familiar. Simplifying the management of the CV stack is especially important for a clear and correct implementation of BLAKE3. This section goes into detail about how this is done in the [reference implementation](#), as an aid to implementers.

5.1.1 Chunk State

The chunk state contains the 32-byte CV of the previous block and a 64-byte input buffer for the next block, and typically also the compression function parameters t and d . Input bytes from the caller are copied into the buffer until it is full. Then the buffer is compressed together with the previous CV, using the truncated compression function. The output CV overwrites the previous CV, and the buffer is cleared. An important detail here is that the last block of a chunk is compressed differently, setting the `CHUNK_END` flag and possibly the `ROOT` flag. In an incremental setting, any block could potentially be the last, until the caller has supplied enough input to place at least 1 byte in the block after. For that reason, the chunk state waits to compress the next block until both the buffer is full and the caller has supplied more input. Note that the CV stack takes the same approach immediately below: chunk CVs are added to the stack only after the caller supplies at least 1 byte for the following chunk.

5.1.2 Chaining Value Stack

To help picture the role of the CV stack, Figure 6 shows a growing tree as chunk CVs are added incrementally. As just discussed above, chunk CVs are added to this tree only after the caller has supplied at least 1 byte for the following chunk, so we know that none of these chunks or parent nodes is the root of the tree, and we do not need to worry about the `ROOT` flag yet.




Figure 6: An incomplete tree growing incrementally from 1 to 4 chunks. Dotted boxes represent CVs that no longer need to be stored.

When the first chunk CV (c_0) is added, it is alone. When the second chunk CV (c_1) is added, it completes a two-chunk subtree, and we can compute the first parent CV (p_0). The third chunk CV (c_2) does not complete any subtrees. Its final position in the tree structure will depend on whether it gets a right sibling (see Figure 1), so we cannot create any parent nodes for it yet. The fourth chunk CV (c_3) provides a right sibling for c_2 and completes two subtrees, one of two chunks and one of four chunks. First it merges with c_2 to compute p_1 , then p_1 merges with p_0 to compute p_2 .

Note that once a subtree is complete, none of its child CVs will be used again, and we do not need to store them any longer. These unneeded CVs are represented by dotted boxes in Figure 6. Solid boxes represent CVs that are still needed. These are what we store in the CV stack, ordered based on when we computed them, with newer CVs on top.

Look through Figure 6 again, this time paying attention to the state of the CV stack at each step. When c_0 is added to the tree, it is the only entry in the stack. When c_1 is added, we pop c_0 off the stack to merge it, and p_0 becomes the only entry. When c_2 is added, there are two entries in the stack, with c_2 on top because it is newer. When c_3 is added, we perform

two merges, first popping off c_2 and then popping off p_0 , and only p_2 remains in the stack.

Note how the stack order of c_2 and p_0 above matched the order in which we needed to retrieve them for merging. By ordering CVs in the stack from newest at the top to oldest at the bottom, we also implicitly order them from tree-right to tree-left, and from the smallest subtree so far to the largest subtree so far. This invariant means that the next CV we need to merge is always on top of the stack.

The key question is then, when a new chunk CV is added to the tree, how many subtrees does it complete? Which is to say, how many CVs should we pop off the stack and merge with it, before pushing on the final result?

To answer this question, note another pattern. The entries in the CV stack behave like the 1 bits in the binary representation of the total number of chunks so far. For example with 3 chunks (0b11), there are two entries in the stack, corresponding to the two 1 bits. With 4 chunks (0b100), only one entry remains, corresponding to the single 1 bit.

To see why this pattern holds, note that all the completed subtrees represented by CVs in the stack contain a number of chunks that is a power of 2. Furthermore, each subtree is a distinct power of 2, because when we have two subtrees of the same size, we always merge them. Thus, the subtrees in the stack correspond to the distinct powers of 2 that sum up to the current total number of chunks. This is equivalent to the 1 bits in the binary representation of that number.

In this sense, popping a CV off the stack to merge it is like flipping a 1 bit to 0, and pushing the final result back onto the stack is like flipping a 0 bit to 1. The fact that we do two merges when adding the fourth chunk, corresponds to the fact that two 1 bits flip to 0 as the total changes from 3 to 4. In other words, we do two merges when we add the fourth chunk, because there are two trailing 1 bits in the number 3, and similarly two trailing 0 bits in the number 4.

This pattern leads to an algorithm for adding a new chunk CV to the tree: For each trailing 0 bit in the new total number of chunks, pop a CV off the stack, and merge it with the new CV we were about to add. Finally, push the fully merged result onto the stack. Listing 1 shows this algorithm as it appears in the Rust reference implementation.

```
fn add_chunk_chaining_value(&mut self, mut new_cv: [u32; 8], mut total_chunks: u64) {  
    // This chunk might complete some subtrees. For each completed subtree,  
    // its left child will be the current top entry in the CV stack, and  
    // its right child will be the current value of `new_cv`. Pop each left  
    // child off the stack, merge it with `new_cv`, and overwrite `new_cv`  
    // with the result. After all these merges, push the final value of  
    // `new_cv` onto the stack. The number of completed subtrees is given  
    // by the number of trailing 0 bits in the new total number of chunks.  
    while total_chunks & 1 == 0 {  
        new_cv = parent_cv(self.pop_stack(), new_cv, self.key, self.flags);  
        total_chunks >>= 1;  
    }  
    self.push_stack(new_cv);  
}
```

Listing 1: The algorithm in the Rust reference implementation that manages the chaining value stack when a new chunk CV is added.

Once the caller indicates that the input is complete, we compute the CV of the current chunk, which may be incomplete or empty. We then merge this CV with each CV in the

current CV stack. This happens regardless of the number of chunks so far, reflecting the fact that subtrees along the right edge of the tree may be incomplete. We set the `ROOT` flag for the last parent node in this process. Or if there are no CVs in the stack, and thus no parents to merge, we set the `ROOT` flag for the chunk.

5.2 Multi-threading

Most of the work of computing a BLAKE3 hash is compressing chunks. Each chunk can be compressed independently, and one approach to multi-threading is to farm out individual chunks or groups of chunks to tasks on a thread pool. In this approach, a leader thread owns the chaining value stack (cf. §5.1.2) and awaits a CV from each task in order.

This leader-workers approach has some inefficiencies. Spawning tasks and creating channels usually require heap allocation, which is a performance cost that needs to be amortized over larger groups of chunks. At high degrees of parallelism, managing the CV stack itself can become a bottleneck.

A more efficient approach to multi-threading is based on a recursive tree traversal. The input is split into left and right parts. As per the rules in §2.1, the left part receives the largest power-of-2 number of chunks that leaves at least 1 byte for the right part. Each part then repeats this splitting step recursively, until the parts are chunk-sized, and each chunk is compressed into a CV. On the way back up the callstack, each pair of left and right child CVs is compressed into a parent CV.

This recursive approach fits well into a fork-join concurrency model, like those provided by OpenMP, Cilk, and Rayon (Rust). Each left and right part becomes a separate task, and a work-stealing runtime parallelizes those tasks across however many threads are available. This can work without heap allocation, because the runtime can make a fixed-size stack allocation at each recursive callsite.

The recursive approach is simplest when the entire input is available at once, since no CV stack is needed. In an incremental setting, a hybrid approach is also possible. A large buffer of input can be compressed recursively into a single subtree CV, and that CV can be pushed onto the CV stack using the same algorithm as in §5.1.2. If each incremental input is a fixed power-of-2 number of chunks in size (for example if all input is copied into an internal buffer before compression), the push algorithm works with no modification. If each input is a variable size (for example if input from the caller is compressed directly without copying), the implementation needs to maintain the largest-to-smallest ordering invariant of the CV stack. The size of each subtree being compressed must evenly divide the total number of input bytes received so far, and the implementation might need to break up the caller's input into separate pieces.

5.3 SIMD

There are two approaches to using SIMD in a BLAKE3 implementation, and both are important for high performance at different input lengths. The first approach is to use 128-bit vectors to represent the 4-word rows of the state matrix. The second approach is to use vectors of any size to represent words in multiple states, which are compressed in parallel.

The first approach is similar to how SIMD is used in BLAKE2b or BLAKE2s, and it is applicable to inputs of any length, particularly short inputs where the second approach does not apply. The state $v_0 \dots v_{15}$ is arranged into four 128-bit vectors. The first vector contains

the state words $v_0 \dots v_3$, the second vector contains the state words $v_4 \dots v_7$, and so on. Implementing the G function (Appendix B) with vector instructions thus mixes all four columns of the state matrix in parallel. A diagonalization step then rotates the words within each row so that each diagonal now lies along a column, and the vectorized G function is repeated to mix diagonals. Finally the state is undiagonalized, to prepare for the column step of the following round.

The second approach is similar to how SIMD is used in BLAKE2bp or BLAKE2sp. In this approach, multiple chunks are compressed in parallel, and each vector contains one word from the state matrix of each chunk. That is, the first vector contains the v_0 word from each state, the second vector contains the v_1 word from each state, and so on, using 16 vectors in total. The width of the vectors determines the number of chunks, so for example 128-bit vectors compress 4 chunks in parallel, and 256-bit vectors compress 8 chunks in parallel. Here the G function operates on one column or diagonal at a time, but across all of the states, and no diagonalization step is required. When enough input is available, this approach is much more efficient than the first approach. It also scales to wider instruction sets like AVX2 and AVX-512.

For SIMD implementations with multiple 128-bit lanes, such as AVX2 and AVX-512, there is a 3rd approach that simply replicates the first approach without transposing the state / message. This might be advantageous in some short-message regimes to be determined.

The second approach can be integrated with the CV stack algorithm from §5.1.2 by computing multiple chunk CVs in parallel and then pushing each of them into the CV stack one at a time. It can also be combined with either of the multi-threading strategies from §5.2. Rather than having each task or recursive leaf compress one chunk at a time, each can compress multiple chunks in parallel.

5.4 Memory Requirements

BLAKE3 has a larger memory requirement than BLAKE2, because of the chaining value stack described in §5.1.2. An incremental implementation needs space for a 32-byte chaining value for every level of the tree below the root. The maximum input size is $2^{64} - 1$ bytes, and the chunk size is 2^{10} bytes, giving a maximum tree depth of $64 - 10 = 54$. The CV stack thus requires $54 \cdot 32 = 1728$ bytes. The chunk state (cf. §5.1.1) also requires at least 104 bytes for the chaining value, the message block, and the chunk counter. The size of the reference implementation is 1880 bytes on the callstack.

For comparison, BLAKE2s has a memory footprint similar to the BLAKE3 chunk state alone, at least 104 bytes. BLAKE2b has twice the chaining value size and block size, requiring at least 200 bytes. And the parallel modes BLAKE2bp and BLAKE2sp both require at least 776 bytes.

Space-constrained implementations of BLAKE3 can save space by restricting the maximum input size. For example, the maximum size of an IPv6 “jumbogram” is $2^{32} - 1$ bytes, or just under 4 GiB. At this size, the tree depth is 22 and the CV stack is $22 \cdot 32 = 704$ bytes. For another example, the maximum size of a TLS record is 2^{14} bytes, or exactly 16 KiB. At this size, the tree depth is 4 and the CV stack is $4 \cdot 32 = 128$ bytes.

6 Applications

As a general-purpose hash function, BLAKE3 is suitable whenever a collision-resistant or preimage-resistant hash function is needed to map some arbitrary-size input to a fixed-length output. BLAKE3 further supports keyed modes—in order to be used as a pseudorandom function, MAC, or key derivation function—as well as streaming and incremental processing features.

6.1 Pseudorandom Function and MAC

Like BLAKE2, BLAKE3 provides a keyed mode, `keyed_hash`. This removes the need for a separate construction like HMAC. The `keyed_hash` mode is also more efficient than keyed BLAKE2 or HMAC for short messages. BLAKE2 requires an extra compression for the key block, and HMAC requires three extra compressions. The `keyed_hash` mode in BLAKE3 does not require any extra compressions.

6.2 Key Derivation

BLAKE3 provides a key derivation mode, `derive_key`. This mode accepts a context string of any length and key material of any length, and it returns a derived key of any length. The context string should be hardcoded, globally unique, and application-specific. A good default format for context strings is "[application] [commit timestamp] [purpose]", e.g., "example.com 2019-12-25 16:18:03 session tokens v1". This mode should not be used with a variable context string, unless that variable is itself a context parameter from a higher level API with the same documented security requirement.

The purpose of this strict requirement is to ensure that there is no way for an attacker in any scenario to cause two different applications or components to inadvertently use the same context string. There should be no dependency whatsoever between user input and a context string used with `derive_key`.

Given this requirement, we then make an explicit exception to an otherwise universal rule of practical cryptography: Applications may use `derive_key` with key material that is already in use with other algorithms. This includes other hash functions, ciphers, and abstractions like HMAC and HKDF. The only limitation in practice is algorithms that forbid key reuse entirely, like a one-time pad. (For other theoretical limitations, see §7.7.)

Normally, using the same key with multiple algorithms is forbidden. Many algorithms are built from the same primitives, and it can be difficult to know when two algorithms might have related output, or when one algorithm might publish a value that another algorithm considers secret. One especially relevant example of this problem is HMAC and HKDF themselves. Although these algorithms have entirely different purposes, HKDF is defined in terms of HMAC. In fact, in the “expand only” mode of HKDF, its output is equivalent to a single call to HMAC. If an application used the same key with HMAC and with HKDF, and an HMAC input happened to collide with an HKDF context string, this interaction could ruin its security.

The standard advice for such an application is to use its original key only with a key derivation function, and to derive independent subkeys for other algorithms. This is good and practical advice in many cases. But it may not be practical when an application or a protocol evolves over time. If a key was originally used only with HMAC, and perhaps years later a second use case arises, the developers may have painted themselves into a corner.

Backwards compatibility might prevent them from using a derived subkey with HMAC. In cases like this, the standard advice amounts to either building a time machine, or over-engineering applications and protocols to derive subkeys at every step, in anticipation of unknown future use cases.

This is why it is valuable for `derive_key` to make an exception to the rule against mixing algorithms. It makes key derivation a practical option not only for new designs, but also for existing applications that are adding new features. In exchange, it is the responsibility of each caller to provide (or for abstractions, to require) a hardcoded, globally unique, application-specific context string, which can never be made to collide with any other.

The `derive_key` mode is intended to replace the BLAKE2 personalization parameter in most of its use cases. Key derivation can encourage better security than personalization, by cryptographically isolating different components of an application from one another. This limits the damage that one component can cause by accidentally leaking its key.

6.3 Stateful Hash Object

To abstract away various ways to use hash function in the Noise protocol framework, Trevor Perrin has suggested the notion of *stateful hash objects* (SHO) [19]. A SHO is an interface with 3 methods:

- `init(label)`: Initialize the internal state using a custom domain-separation label.
- `absorb(data)`: Hashes an arbitrary number of bytes into the state.
- `squeeze(length)`: Produces an arbitrary amount of pseudorandom output based on the current state.

Due to its zero-overhead keying and variable-length output, BLAKE3 may be adapted into a SHO quite easily. The internal state at the time t is represented by a 256-bit key k_t :

- The initial key, k_0 , is obtained as $k_0 = \text{hash}(\text{label})$;
- `absorb(data)` updates the key: $k_{i+1} = \text{keyed_hash}(k_i, \text{data})$;
- `squeeze(length)` is done naturally by extending the output of the previous `absorb(data)` to more than 32 bytes¹.

6.4 Verified Streaming

Because BLAKE3 is a tree hash, it supports new use cases that serial hash functions do not. One new use case is verified streaming. Consider a video streaming application that fetches video files from an untrusted host. The application knows the hash of the file it wants, and it needs to verify that the video data it receives matches that hash. With a serial hash function, nothing can be verified until the application has downloaded the entire file. But with BLAKE3, verified streaming is possible. The application can verify and play individual chunks of video as soon as they arrive.

¹Extending the output of the last `absorb` entails preserving the last up to 64 bytes of input data in memory for possibly a long time; in such cases it might be desirable to call `absorb(empty)` after each `absorb` call.

To verify an individual chunk without re-hashing the entire file, we verify each parent node on the path from the root to that chunk. For example, suppose the file is composed of four chunks, like the four-chunk tree in Figure 1. To verify the first chunk, we start by fetching the root node. (Specifically, we fetch its message bytes, the concatenated chaining values of its children.) We compute the root node's CV as per §2.5 and confirm that it matches the 32-byte BLAKE3 hash of the entire file. Then, we fetch the root node's left child, which is the first chunk's parent. We compute that node's CV and confirm that it matches the first 32-bytes of the root node. Finally, we fetch the first chunk itself. We compute its CV as per §2.4 and verify that it matches the first 32-bytes of its parent. This verifies that the first chunk is authentic, and we can pass it along to application code.

To continue streaming, we can immediately fetch the second chunk. It shares the parent node of the first chunk, and its CV should match the second 32 bytes of that parent node. For the third chunk, we need to fetch its parent. The CV of that parent node should match the second 32 bytes of the root node, and then the CV of the third chunk should match the first 32 bytes of its parent. Note that whenever we fetch a parent node, we immediately use its first 32 bytes to check its left child's CV, and then we store its second 32 bytes to check its right child in the future. We can represent this with a stack of expected CVs. We push CVs onto this stack and pop them off, as we would with node pointers in a depth-first traversal of a binary tree. Note that this is different from the "CV stack" used for incremental hashing in §5.1.2; that stack holds CVs we have computed in the past, while this stack holds CVs we expect to compute in the future, and we manage them with different algorithms.

Observe that if the application eventually verifies the entire file, it will have fetched all the nodes of the tree in pre-order. This suggests a simple wire format for a streaming protocol: The host can concatenate all the parent nodes and chunks together in pre-order and serve the concatenated bytes as a single stream. If the client application knows the length of the file in advance, it does not need any other information to parse the stream and verify each node. In this format, the space overhead of the parent nodes approaches two CVs per chunk, or 6.25%.

If the client does not know the length of the file in advance, it may receive the length from the host, either separately or at the front of the stream. In this case, the length is untrusted. If the host reports an incorrect length, the effect will be that at least the final chunk will fail verification. For this reason, if an application reads the file sequentially, it does not need to explicitly verify the length. The stream may terminate with an error partway through if verification fails, as it might if the network connection failed. An important edge case is that, if the reported length is 0, the file consists of a single empty chunk, and the implementation must not forget to verify that the file hash matches the empty chunk's CV. (One way to make this mistake is for the implementation to return end-of-file as soon as the current read position equals the expected length, verifying nothing in the zero-length case.)

On the other hand, if an application implements seeking, length verification is required. Seeking past the reported end-of-file, or performing an EOF-relative seek, would mean trusting an unverified length. In these cases, the implementation must first verify the length by seeking to and verifying the final chunk. If that succeeds, the length is authentic, and the implementation can proceed with the caller's requested seek.

This verified streaming protocol has been implemented by the Bao tool. It is conceptually similar to the existing Tree Hash Exchange format, and to the BEP 30 extension of the BitTorrent protocol. However, neither of those prevents length extension, and the latter (if extracted from the BitTorrent protocol) does not provide second-preimage resistance [11,

§8.5].

6.5 Incremental Update

Another new use case supported by BLAKE3 is incrementally updating the root hash when an input is edited in place. A serial hash function can efficiently append new bytes to the end of an input, but editing bytes at the beginning or in the middle requires re-hashing everything to the right of the edit. With BLAKE3, an application can edit bytes anywhere in its input and re-hash just the modified chunks and their transitive parent nodes.

For example, consider an input composed of four chunks, again like the four-chunk tree in Figure 1. Suppose we overwrite some bytes in the second chunk. To update the root hash, we first compute the second chunk's new CV as per §2.4. Then we recompute the CV for the second chunk's parent as per §2.5. Note that because the first chunk is unchanged, the first 32 message bytes for this parent node are unchanged also, and we do not need to re-read the first chunk. Finally, we recompute the CV of the root node, and we similarly do not need to re-read anything from the right half of the tree.

Note that this strategy cannot efficiently insert or remove bytes from the middle of an input. That would change the position of all the bytes to the right of the edit and force re-hashing. This is similar to the constraints of a typical filesystem, where appends and in-place edits are efficient, but insertions and removals within a file are either slow or unsupported.

7 Rationales

This section goes into more detail about some of the performance tradeoffs and security considerations in the BLAKE3 design.

7.1 Chunk Size

BLAKE3 uses 1 KiB chunks. The chunk size is a tradeoff between the peak throughput for long inputs, which benefits from larger chunks, and the degree of parallelism for medium-length inputs, which benefits from shorter chunks.

The benefit of a larger chunk size is that the tree contains fewer parent nodes, so the overhead cost of compressing them is lower. The number of parent nodes is equal to the number of chunks minus one, so doubling the chunk size cuts the number of parent nodes in half. However, note that parent nodes at the same level of the tree can be compressed in parallel, so having twice as many parent nodes does not necessarily mean that compressing them takes twice as long.

The benefit of a smaller chunk size is a higher potential degree of parallelism. This is irrelevant for very long inputs, where all the SIMD lanes and possibly CPU cores of the machine will be occupied regardless. But at medium input lengths, this has a huge impact on performance. Consider an 8 KiB input. With a chunk size of 1 KiB, this input can occupy 8 SIMD lanes, which lets the implementation use AVX2 vector instructions on modern x86 platforms. But if we increased the chunk size to 2 KiB, the same input would only occupy 4 SIMD lanes, and the high throughput of AVX2 would be left on the table.

Thus we want to pick the smallest chunks size we can to maximize medium-length parallelism, without incurring “excessive” overhead from parent nodes. The point of comparison

here is the peak throughput for a very large chunk size, like 64 KiB, where parent node overhead becomes small enough to be negligible. In our measurements, we find that implementations with 1 KiB chunks can reach at least 90% of that peak throughput on both modern x86 platforms and on the ARM1176.

There are other minor performance considerations to be aware of. A larger chunk size slightly reduces the space requirement of the CV stack as described in §5.4. Also, verified streaming requires buffering chunks (see §6.4), and incremental updates require re-hashing chunks (see §6.5), so both of those use cases benefit from a smaller chunk size.

7.2 Word Size

BLAKE3 uses 32-bit words. The performance tradeoffs in the choice between 32-bit and 64-bit words are complex, and they depend on both the platform and the size of the input.

Many sources note that SHA-512 (because it uses 64-bit words) is faster on 64-bit platforms than SHA-256 (because it uses 32-bit words). A similar effect applies to BLAKE2b and BLAKE2s. However, this effect does not apply to algorithms that incorporate more SIMD parallelism. BLAKE2sp has higher peak throughput than BLAKE2bp on x86-64. This is because the SIMD instructions that operate on 8 lanes of 32-bit words have the same throughput as those that operate on 4 lanes of 64-bit words. The key factor for exploiting SIMD performance is not word size but vector size, and both of those cases occupy a 256-bit vector. (The remaining effect making BLAKE2sp faster is that its smaller state requires fewer rounds of compression.)

For that reason, peak throughput on x86-64 is largely independent of the word size. The performance differences that remain are restricted to shorter input lengths: 32-bit words perform better for lengths less than one block, while 64-bit words perform better for lengths between one block and one chunk, comparable to BLAKE2s and BLAKE2b. These effects are minor, and short input performance is usually dominated by other sources of overhead in typical applications.

Instead, the decisive advantage of 32-bit words over 64-bit words is performance on smaller architectures and embedded systems. BLAKE3 aims to be a single general-purpose hash function, suitable for users of both BLAKE2b and BLAKE2s. On 32-bit platforms like the ARM1176, where BLAKE2s does especially well, a 64-bit hash function could be a performance regression. Choosing 32-bit words is a substantial benefit for these platforms, with little or no downside for x86-64. Cross-platform flexibility is also important for protocol designers, who have to consider the performance of their designs across a wide range of hardware, and who are becoming increasingly skeptical of designs supporting multiple algorithms. [13]

7.3 Compression Function Feed Forward

The output step of the compression function (see §4.3) mixes the first and second halves of the state, and then feeds forward the input CV into the second half. The feed-forward step allows us to use all 16 words of the state as output, doubling the output rate of the extendable output feature (see §2.6). If we exposed all 16 state words without the feed-forward, the compression function would be reversible when the message bytes were known.

When the full output is used, mixing the first and second halves of the state is redundant. However, when the truncated output is used, mixing the halves makes the compression

function slightly stronger against truncated differentials. The default 32-byte output size is a truncated output, and truncated outputs are also used for chaining values in the interior of the tree.

7.4 Chunk Counter

The main job of the compression function's 64-bit counter parameter, t , is to support extendable output. By incrementing t (see §2.6), the caller can extract as much output as needed, without imposing any extra steps on the default-length case.

However, we also use the t parameter during the input phase, to indicate the chunk number (see §2.4). This means that each chunk has a distinct CV with high probability, even if two chunks have the same input bytes. This is not strictly necessary for security, but it discourages a class of dangerous optimizations.

Consider an application that hashes sparse files, which are mostly filled with zeros. The majority of this application's input chunks are the all-zero chunk. This application might try to compute the CV of the zero chunk only once, and then before compressing each chunk of input, check to see whether its bytes are all zero. This check is cheap, and for this application it could be an enormous speedup. But this optimization is dangerous, because it is not constant-time. The runtime of the hash function would leak information about its input.

If tricks like this were possible, an unsafe implementation would inevitably find its way into some privacy-sensitive use case. By distinguishing each chunk, BLAKE3 deliberately sabotages these tricks, with the hope of keeping every implementation constant-time.

7.5 Tree Fanout

BLAKE3 uses a binary tree structure, which is to say a fanout of 2. A binary tree is simpler to implement than a wider tree, and its memory requirement is also a local minimum.

The CV stack algorithm described in §5.1.2 is important for simplifying a BLAKE3 implementation. It also saves space, because the implementation does not need to store size or level metadata along with each CV. Consider how this algorithm would change, if BLAKE3 used a fanout of 4. Instead of "count the trailing 0 bits, and pop that many CVs" it would be "count the trailing 0-0 bit pairs, and pop that many CV triplets." That can still be done with a bit mask, but programmers would need to reason about quaternary numbers to understand what was going on. Other complexities would crop up as well: parent nodes would no longer be a fixed size, and new ambiguities in the tree layout would require a third rule in §2.1 along the lines of "all non-rightmost sibling subtrees have the same number of chunks."

A fanout of 4 would also increase the memory requirement of the CV stack. Instead of storing one CV per level of the tree, we would need to store three. The tree would be half as tall, but ultimately the CV stack would be larger by a factor of $3/2$. In general, the size of the CV stack for a fanout F is proportional to $(F - 1)/\log_2(F)$, which increases as F increases.

A different approach to the CV stack is possible, which could save space at larger fanouts. Rather than storing CV bytes directly, the implementation could store an incremental state for each incomplete parent node, similar to the incremental chunk state. An incremental state requires a 32-byte CV and a 64-byte buffer, so at fanout 4 it would be no better than storing 3 CVs directly. (Perhaps the 64-byte buffer could be reduced to 32 bytes in some cases, depending on the details of parent node finalization, and at the cost of more complexity.) But

at higher fanouts, the memory requirement would be proportional to $1/\log_2(F)$. At fanout 16 and above, this approach could save space relative to the binary tree.

However, recall from §5.4 that space-constrained implementations of BLAKE3 can reduce the size of the CV stack by restricting their maximum input length. We expect that the majority of space-constrained applications already have limits on the size of their input. Indeed it is hard to imagine how an application without 2 KB of memory to spare could get its hands on a gigabyte of input. Realistically, a large fanout would only benefit a tiny fringe of space-constrained applications, but the added complexity would hurt everyone.

As a special case, we could also have chosen an infinite fanout, a single parent node with an unlimited number of leaves. This approach does not require a CV stack at all. Instead, the implementation only needs to manage the incremental state of the current chunk and the incremental state of the single parent node. Notably, this is the approach taken by the KangarooTwelve hash function. It has large benefits for simplicity and space efficiency, but also several downsides. Parent blocks cannot be compressed in parallel, requiring a larger chunk size to keep overhead low (see §7.1). Multi-threading is more complex and generally requires heap allocation, because the input cannot be split in half recursively (see §5.2). And incremental use cases like verified streaming (see §6.4) are no longer practical.

As a totally different approach, we could also have chosen a fixed-interleaved tree layout similar to that of BLAKE2bp and BLAKE2sp. For example, BLAKE2sp divides its input into 8 interleaved pieces, with blocks 0, 8, 16... forming the first piece, blocks 1, 9, 17... forming the second piece, and so on. The 8 resulting CVs are then concatenated into a single parent node. The upsides of this approach are that full parallelism can begin at relatively short input lengths, 512 bytes in the case of BLAKE2bp and BLAKE2sp, and that the memory requirement is lower than that of the binary tree. The downsides of this approach are that the degree of parallelism is fixed, and that performance is poor for shorter inputs.

7.6 Domain Flags

The `ROOT` and `CHUNK_START` domain flags are necessary for security. The `ROOT` flag separates the root node from all other nodes, and it separates the final block from all other blocks in a root chunk, both of which prevent length extensions. The `CHUNK_START` flag separates chunks from parent nodes, and it separates an IV provided by the caller (the key in the keyed hashing mode) from block CVs computed within a chunk, both of which prevent collisions.

However, the `PARENT` and `CHUNK_END` domain flags are not strictly necessary for security, and we include them to be conservative. As with `CHUNK_START`, the `PARENT` flag prevents a collision when the caller-provided IV bytes match a block CV extracted from some other input. However, the parent's message bytes will also be different from the block's with high probability in that case. The benefit of the `PARENT` flag is rather that we get separation without needing to reason about the message bytes.

The `CHUNK_END` flag prevents length extensions for all chunks, beyond just a root chunk. However, other chunk CVs are not usually published. In incremental use cases like verified streaming (see §6.4), chunk CVs may be published, but only when the input is also published. In a keyed incremental use case, the keyed parent nodes also prevent an attacker from making any use of length-extended keyed chunk CVs. Nonetheless, we include the `CHUNK_END` flag, because letting an attacker length-extend chunk CVs is an unnecessary risk that might impact unknown future applications.

7.7 Key Derivation from Reused Key Material

As described in §6.2, the `derive_key` mode is intended to be safe to use even with keys that are already in use with other algorithms. It is impossible to guarantee such a property in general, however. As noted in that section, some algorithms explicitly forbid key reuse, even within the same algorithm. In the case of a one-time pad, if the message happens to be known or predictable, some or all of the key is revealed to an eavesdropper. No key derivation function can extract useful security from key material that is not secret.

There are some algorithms that impose a one-time-use restriction on their key, but that still keep the key secret as long as that restriction is followed. The `crypto_onetimeauth` function in the NaCl and libsodium libraries is one such example. It might be safe to use the same key with `crypto_onetimeauth` (once) and with `derive_key` (any number of times). However, precisely because of its one-time-use restriction, `crypto_onetimeauth` is almost always used with derived subkeys anyway. Furthermore, this sort of algorithm-specific reasoning is exactly what `derive_key` is designed to avoid. It is simpler and safer to forbid `derive_key` from sharing key material with this entire class of one-time-use algorithms.

Another limitation on key reuse is applications that violate the security requirement of the `derive_key` context string, namely that it should be hardcoded, globally unique, and application-specific. If one component of an application feeds arbitrary user input into the context string, it is not safe for any other component to share key material with it. The context string no longer provides domain separation. This is the basic problem of sharing key material, and the reason that deriving separate subkeys is preferable wherever possible. When two different components share a key, each has to assume that the other will not violate the security requirements of `derive_key`, just as each has to assume that the other will not leak the key.

Another more theoretical limitation for key material reuse could be future algorithms that are internally identical to BLAKE3, but with different domain-separation conventions. A modified algorithm might sabotage domain separation entirely, for example, by exposing the `d` parameter of the compression function to arbitrary user input. We will never publish such an algorithm. The common practice when designing a closely related algorithm is to change the IV constants, so that the new output is independent. This was the approach used to derive SHA-512/256 from SHA-512. The same approach could work for a hypothetical derivative of BLAKE3 by changing the $IV_0 \dots IV_3$ constants, which are used for compression function setup in all modes.

Acknowledgments

Thanks to Liming Luo for many discussions about the tree layout.

References

- [1] Kevin Atighehchi and Alexis Bonnetaze. Asymptotic analysis of plausible tree hash modes for SHA-3. *IACR Trans. Symmetric Cryptol.*, 2017(4):212–239, 2017. doi: [10.13154/tosc.v2017.i4.212-239](https://doi.org/10.13154/tosc.v2017.i4.212-239).

- [2] Kevin Atighehchi and Robert Rolland. Optimization of tree modes for parallel hash functions: A case study. *IEEE Trans. Computers*, 66(9):1585–1598, 2017. doi:[10.1109/TC.2017.2693185](https://doi.org/10.1109/TC.2017.2693185).
- [3] Jean-Philippe Aumasson. Too much crypto. In *Real World Crypto*, 2020.
- [4] Jean-Philippe Aumasson, Jian Guo, Simon Knellwolf, Krystian Matusiewicz, and Willi Meier. Differential and invertibility properties of BLAKE. In Seokhie Hong and Tetsu Iwata, editors, *Fast Software Encryption, 17th International Workshop, FSE 2010, Seoul, Korea, February 7-10, 2010, Revised Selected Papers*, volume 6147 of *Lecture Notes in Computer Science*, pages 318–332. Springer, 2010. doi:[10.1007/978-3-642-13858-4_18](https://doi.org/10.1007/978-3-642-13858-4_18).
- [5] Jean-Philippe Aumasson, Samuel Neves, Zooko Wilcox-O’Hearn, and Christian Winerlein. BLAKE2: simpler, smaller, fast as MD5. In Michael J. Jacobson Jr., Michael E. Locasto, Payman Mohassel, and Reihaneh Safavi-Naini, editors, *Applied Cryptography and Network Security - 11th International Conference, ACNS 2013, Banff, AB, Canada, June 25-28, 2013. Proceedings*, volume 7954 of *Lecture Notes in Computer Science*, pages 119–135. Springer, 2013. doi:[10.1007/978-3-642-38980-1_8](https://doi.org/10.1007/978-3-642-38980-1_8).
- [6] Dongxia Bai, Hongbo Yu, Gaoli Wang, and Xiaoyun Wang. Improved boomerang attacks on round-reduced SM3 and keyed permutation of BLAKE-256. *IET Information Security*, 9(3):167–178, 2015. doi:[10.1049/iet-ifs.2013.0380](https://doi.org/10.1049/iet-ifs.2013.0380).
- [7] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Sufficient conditions for sound tree and sequential hashing modes. *Int. J. Inf. Sec.*, 13(4):335–353, 2014. doi:[10.1007/s10207-013-0220-y](https://doi.org/10.1007/s10207-013-0220-y).
- [8] Alex Biryukov, Daniel Dinu, and Dmitry Khovratovich. Argon2: New generation of memory-hard functions for password hashing and other applications. In *IEEE European Symposium on Security and Privacy, EuroS&P 2016, Saarbrücken, Germany, March 21-24, 2016*, pages 292–302. IEEE, 2016. doi:[10.1109/EuroSP.2016.31](https://doi.org/10.1109/EuroSP.2016.31).
- [9] Alex Biryukov, Ivica Nikolic, and Arnab Roy. Boomerang attacks on BLAKE-32. In Antoine Joux, editor, *Fast Software Encryption - 18th International Workshop, FSE 2011, Lyngby, Denmark, February 13-16, 2011, Revised Selected Papers*, volume 6733 of *Lecture Notes in Computer Science*, pages 218–237. Springer, 2011. doi:[10.1007/978-3-642-21702-9_13](https://doi.org/10.1007/978-3-642-21702-9_13).
- [10] Joan Daemen, Tony Dusenge, and Gilles Van Assche. Sufficient conditions for sound hashing using a truncated permutation. *IACR Cryptology ePrint Archive*, 2011:459, 2011. URL: <http://eprint.iacr.org/2011/459>.
- [11] Joan Daemen, Bart Mennink, and Gilles Van Assche. Sound hashing modes of arbitrary functions, permutations, and block ciphers. *IACR Trans. Symmetric Cryptol.*, 2018(4):197–228, 2018. URL: <https://www.cs.ru.nl/~bmennink/pubs/18fsenew.pdf>.
- [12] Yevgeniy Dodis, Leonid Reyzin, Ronald L. Rivest, and Emily Shen. Indifferentiability of permutation-based compression functions and tree-based modes of operation, with applications to MD6. In Orr Dunkelman, editor, *Fast Software Encryption, 16th International Workshop, FSE 2009, Leuven, Belgium, February 22-25, 2009, Revised Selected*

- Papers*, volume 5665 of *Lecture Notes in Computer Science*, pages 104–121. Springer, 2009. doi:10.1007/978-3-642-03317-9_7.
- [13] Jason A. Donenfeld. Wireguard: Next generation kernel network tunnel. In *Proceedings of the Network and Distributed System Security Symposium, NDSS 2017*, 2015. URL: <https://www.wireguard.com/papers/wireguard.pdf>.
 - [14] Orr Dunkelman and Dmitry Khovratovich. Iterative differentials, symmetries, and message modification in BLAKE-256. In *ECRYPT2 Hash Workshop*, 2011.
 - [15] Thomas Espitau, Pierre-Alain Fouque, and Pierre Karpman. Higher-order differential meet-in-the-middle preimage attacks on SHA-1 and BLAKE. In Rosario Gennaro and Matthew Robshaw, editors, *Advances in Cryptology - CRYPTO 2015 - 35th Annual Cryptology Conference, Santa Barbara, CA, USA, August 16-20, 2015, Proceedings, Part I*, volume 9215 of *Lecture Notes in Computer Science*, pages 683–701. Springer, 2015. doi:10.1007/978-3-662-47989-6_33.
 - [16] Jian Guo, Pierre Karpman, Ivica Nikolic, Lei Wang, and Shuang Wu. Analysis of BLAKE2. In Josh Benaloh, editor, *Topics in Cryptology - CT-RSA 2014 - The Cryptographer's Track at the RSA Conference 2014, San Francisco, CA, USA, February 25-28, 2014. Proceedings*, volume 8366 of *Lecture Notes in Computer Science*, pages 402–423. Springer, 2014. doi:10.1007/978-3-319-04852-9_21.
 - [17] Yonglin Hao. The boomerang attacks on BLAKE and BLAKE2. In Dongdai Lin, Moti Yung, and Jianying Zhou, editors, *Information Security and Cryptology - 10th International Conference, Inscrypt 2014, Beijing, China, December 13-15, 2014, Revised Selected Papers*, volume 8957 of *Lecture Notes in Computer Science*, pages 286–310. Springer, 2014. doi:10.1007/978-3-319-16745-9_16.
 - [18] Ji Li and Liangyu Xu. Attacks on round-reduced BLAKE. *IACR Cryptology ePrint Archive*, 2009:238, 2009. URL: <http://eprint.iacr.org/2009/238>.
 - [19] Trevor Perrin. Stateful hash objects: API and constructions, 2018. URL: https://github.com/noiseprotocol/sho_spec.
 - [20] Bozhan Su, Wenling Wu, Shuang Wu, and Le Dong. Near-collisions on the reduced-round compression functions of Skein and BLAKE. In Swee-Huay Heng, Rebecca N. Wright, and Bok-Min Goi, editors, *Cryptology and Network Security - 9th International Conference, CANS 2010, Kuala Lumpur, Malaysia, December 12-14, 2010. Proceedings*, volume 6467 of *Lecture Notes in Computer Science*, pages 124–139. Springer, 2010. doi:10.1007/978-3-642-17619-7_10.
 - [21] Janos Vidali, Peter Nose, and Enes Pasalic. Collisions for variants of the BLAKE hash function. *Inf. Process. Lett.*, 110(14-15):585–590, 2010. doi:10.1016/j.ipl.2010.05.007.

Appendix A IV Constants

The constants $IV_0 \dots IV_7$ used by the compression function are the same as in BLAKE2s. They are:

$$\begin{array}{ll} IV_0 = 0x6a09e667 & IV_1 = 0xbb67ae85 \\ IV_2 = 0x3c6ef372 & IV_3 = 0xa54ff53a \\ IV_4 = 0x510e527f & IV_5 = 0x9b05688c \\ IV_6 = 0x1f83d9ab & IV_7 = 0x5be0cd19 \end{array}$$

Appendix B Round Function

The compression function transforms the internal state $v_0 \dots v_{15}$ through a sequence of 7 rounds. The round function is the same as in BLAKE2s. A round does:

$$\begin{array}{llll} G_0(v_0, v_4, v_8, v_{12}) & G_1(v_1, v_5, v_9, v_{13}) & G_2(v_2, v_6, v_{10}, v_{14}) & G_3(v_3, v_7, v_{11}, v_{15}) \\ G_4(v_0, v_5, v_{10}, v_{15}) & G_5(v_1, v_6, v_{11}, v_{12}) & G_6(v_2, v_7, v_8, v_{13}) & G_7(v_3, v_4, v_9, v_{14}) \end{array}$$

That is, a round applies a G function to each of the columns in parallel, and then to each of the diagonals in parallel. $G_i(a, b, c, d)$ is defined as follows. \oplus denotes bitwise exclusive-or, \ggg denotes bitwise right rotation, and $m_{\sigma_r(x)}$ is the message word whose index is the x^{th} entry in the message schedule for round r :

$$\begin{array}{l} a \leftarrow a + b + m_{\sigma_r(2i)} \\ d \leftarrow (d \oplus a) \ggg 16 \\ c \leftarrow c + d \\ b \leftarrow (b \oplus c) \ggg 12 \\ a \leftarrow a + b + m_{\sigma_r(2i+1)} \\ d \leftarrow (d \oplus a) \ggg 8 \\ c \leftarrow c + d \\ b \leftarrow (b \oplus c) \ggg 7 \end{array}$$

The message schedules σ_r are:

σ_0	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
σ_1	14	10	4	8	9	15	13	6	1	12	0	2	11	7	5	3
σ_2	11	8	12	0	5	2	15	13	10	14	3	6	7	1	9	4
σ_3	7	9	3	1	13	12	11	14	2	6	5	10	4	0	15	8
σ_4	9	0	5	7	2	4	10	15	14	1	11	12	6	8	3	13
σ_5	2	12	6	10	0	11	8	3	4	13	7	5	15	14	1	9
σ_6	12	5	1	15	14	13	4	10	0	7	6	3	9	2	8	11